a Frozen per-exam features and adjacent changes

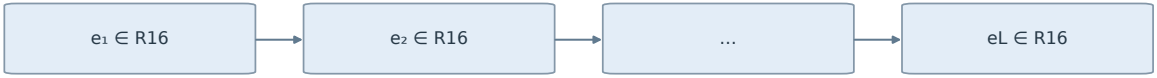

Each e: AFTER ExamSetEncoder, Linear14→32 / LN / GELU / Linear32→16 / set mean.

Only pre-query REAL exams enter inference; Δk = e(k+1) − ek.

b Shared delta encoder and valid-mask pooling

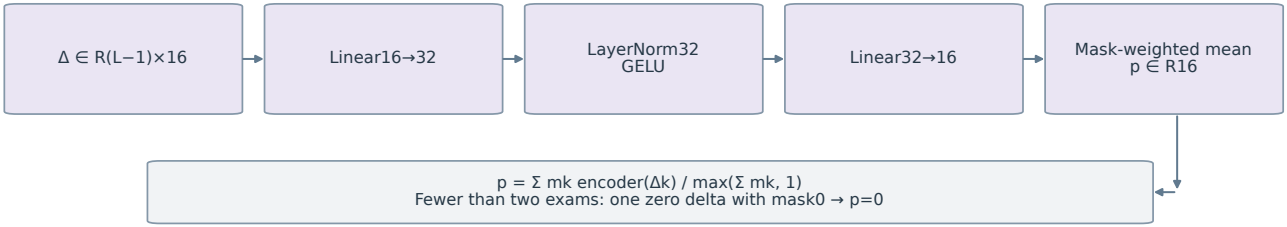

c Four-input fusion and the separate output adapter

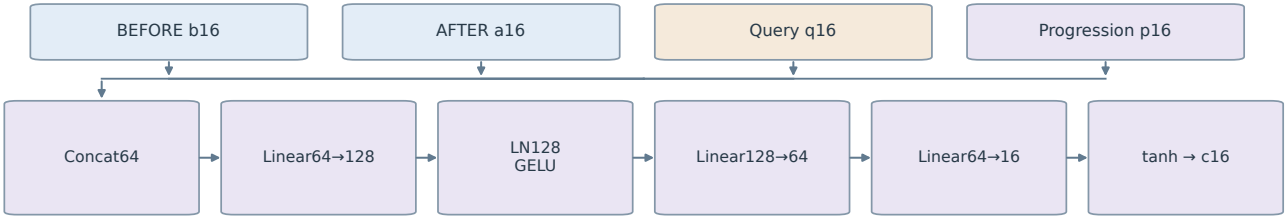

Fusion hidden64; cpfm_adapter is Linear64→16, NOT the B14 controller. Parameters: 19,008.

d Original historical supervision, not new training